week11

- revisit RL basic of spinning tutorial for the list of Key Papers in DRL and updates for the new knowledge
- Step into the next step: after get a basic (relative deep than before) understanding of RL, this time to "actually" get into DRL. The next step is recommended on the website – First get into the field of DL

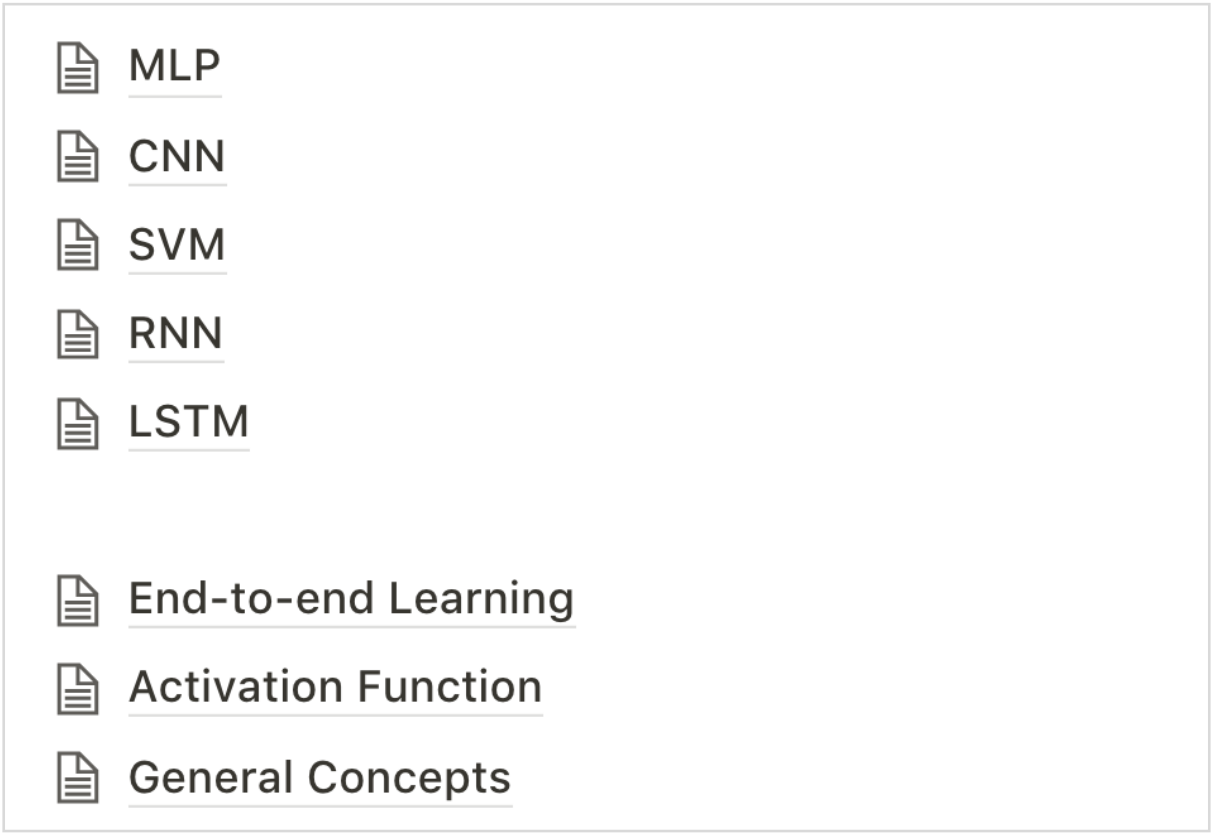

- Begin the second online course: CS294-112

Lecture 1: Introduction and Course Overview
Lecture 2: Supervised Learning and Imitation
Lecture 3: TensorFlow and Neural Nets Review Session (notebook)
Lecture 4: Reinforcement Learning Introduction
Lecture 5: Policy Gradients Introduction
Lecture 6: Actor-Critic Introduction
Lecture 7: Value Functions and Q-Learning
Lecture 8: Advanced Q-Learning Algorithms
Lecture 9: Advanced Policy Gradients
Lecture 10: Optimal Control and Planning
Lecture 11: Model-Based Reinforcement Learning
Lecture 12: Advanced Model Learning and Images
Lecture 13: Learning Policies by Imitating Other Policies
Lecture 14: Probability and Variational Inference Primer
Lecture 15: Connection between Inference and Control
Lecture 16: Inverse Reinforcement Learning
Lecture 17: Exploration: Part 1
Lecture 18: Exploration: Part 2
Lecture 19: Transfer Learning and Multi-Task Learning
Lecture 20: Meta-Learning
Lecture 21: Parallelism and RL System Design
Lecture 22: Advanced Imitation Learning and Open Problems
Lecture 23: Guest Lecture: Craig Boutilier
Lecture 24: Guest Lecture: Gregory Kahn
Lecture 25: Guest Lecture: Quoc Le & Barret Zoph
Lecture 26: Guest Lecture: Karol Hausman (Canceled)
Lecture 27: Final Project Presentations: Part 1 (No Slides)
Lecture 28: Final Project Presentations: Part 2 (No Slides)

- Set up plan to read more paper in the following several weeks

- Rearrange the repo to record the EXACT learning process more accurately than before